Sebastián Pastor Ferrari

Data Mining – Assignment 01

02.03.2020

Question 1 – Data Exploration :

a) **Generate a random sample from the data set that contains 80% of the rows (use the seed = 4279), and answer the following questions with respect to this sample.**
Done.

b) **How many rows are in this data set? How many columns? What do the rows and columns represent?**
There are 506 rows, which each represents a different household in the Boston region. The columns are all different characteristics for each household.

c) **Make some pairwise scatterplots of the predictors (columns) in this data set. Describe your findings.**

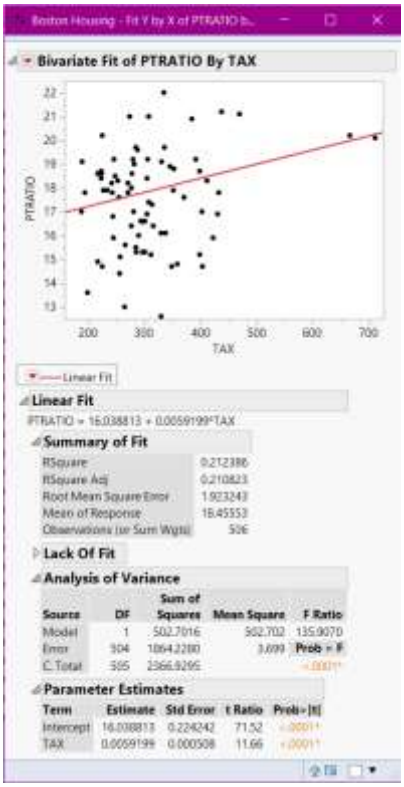

i. x(property tax) by y(pupils per teacher). Property tax is not a good indicator of pupils per teacher. I would have expected that towns with more property tax revenue, to have larger investments in schools (by number of teachers). The opposite is true however, as property taxes go up, the number of pupils-to-teacher goes up as well.

ii. x(weighted distance to five employment centers) by y(proportion of households built before 1940). I expected newer houses to be further away from dense employment centers due to the growth of suburban towns. My assessment was correct in this case.

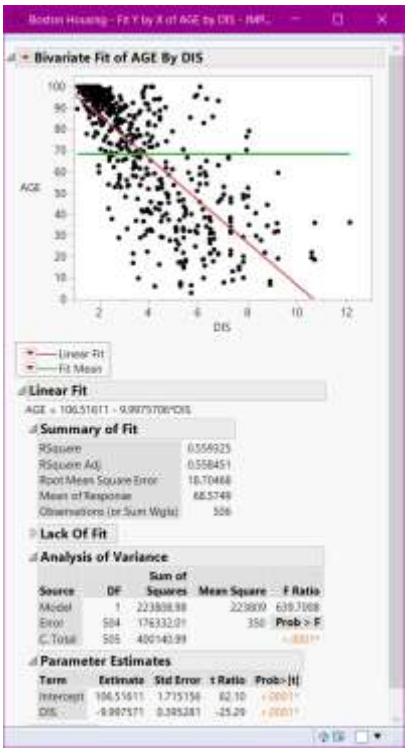

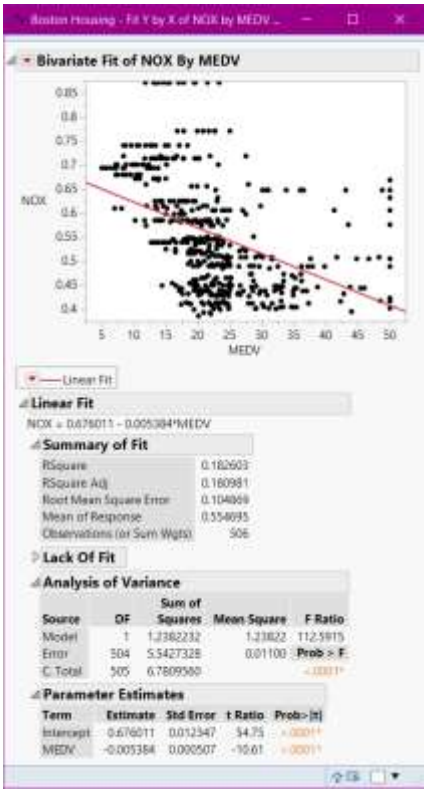

iii. x(Median Value of house) by y(Nitric Oxide Concentration). I expected the concentration of NO in the environment to go down in more expensive towns. However, there seems to be no pattern. Either the people of Boston are not aware of pollution or they don't seem to care about it, indicated by the housing market data we have.

d) Are any of the predictors associated with per capita crime rate? If so, explain the relationship.

Correlations

	CRIM	ZN	INDUS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	LSTAT	MEDV
CRIM	1.0000	-0.2005	0.4066	0.4210	-0.2192	0.3527	-0.3797	0.6255	0.5828	0.2899	0.4556	-0.3883

Per Capita Crime rate is associated with both RAD(Index of Accessibility to highways), and TAX(Full-Value Property Tax Rate). CRIM/RAD makes sense due to criminals wanting good ways of getting away and probably perform crimes near gas stations, fast foods, etc. Full property tax value is surprising, the reasonable explanation for this relationship could be that homes in the metropolitan have higher taxes, and crime takes more place in urban regions.

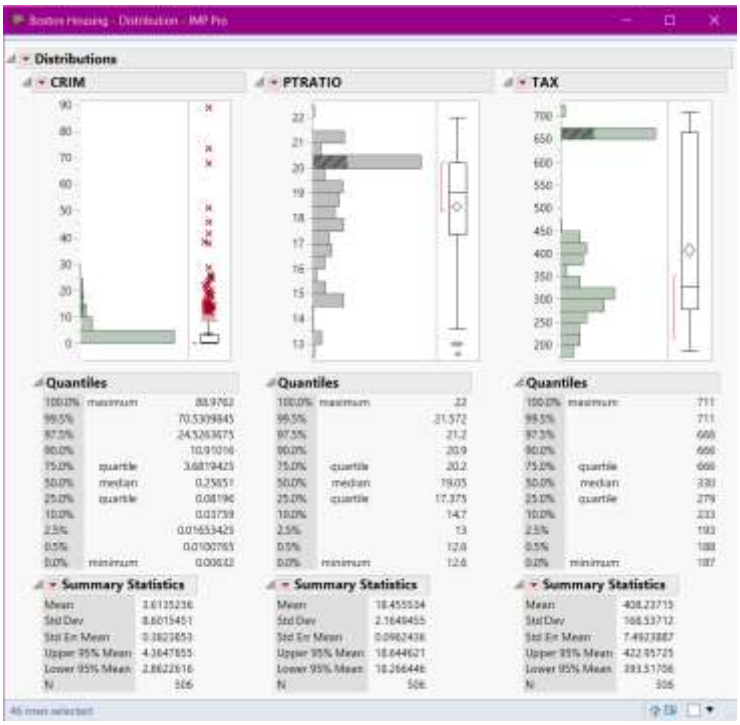

e) Do any of the suburbs of Boston appear to have particularly high crime rates? Tax rates? Pupil-teacher ratios? Comment on the range of each predictor.

There is a range of suburbs that are not marginally, but greatly away from the normal range, particularly in Crime Rate. The mean value for crime rate is 3.6, yet a great deal of towns go from 10 all the way to 90. These towns have particularly high pupil to teacher ratio and high property tax ratios. The Pupil to teacher Ratio has a mean of 18.45, and Tax has a ratio of 408.23.

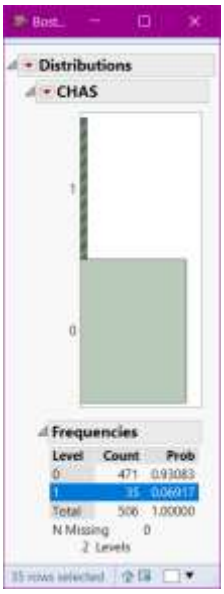

f) How many of the suburbs in this data set bound the Charles river?

35 out of the 506 towns are set bound the Charles River.

g) What is the median pupil-teacher ratio among the towns in this data set?

As mentioned above, the mean of Pupil to Teacher ratio is 18.45.

h) Which suburb of Boston has lowest median value of owner occupied homes? What are the values of the other predictors for that suburb, and how do those values compare to the overall ranges for those predictors? Comment on your findings.

Rows 309 and 406 both have the minimum value (5). They are characterized by high Age, Tax, Crime Rate, and Nitrous Oxide levels.

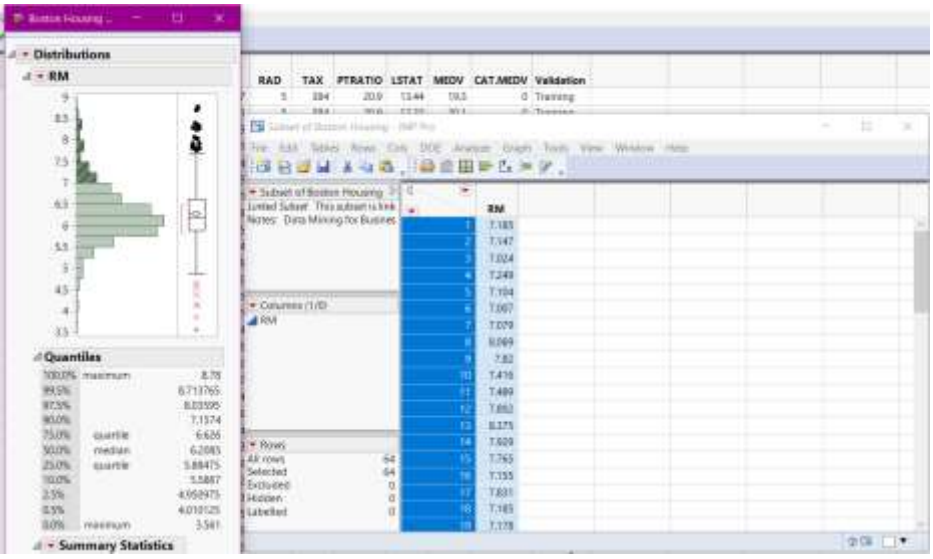

i) In this data set, how many of the suburbs average more than seven rooms per dwelling? More than eight rooms per dwelling? Comment on the suburbs that average more than eight rooms per dwelling.

64 suburbs average more than 7 rooms per dwelling.

13 average more than 8 rooms per dwelling. These have high Median Values, Large Size, Large Age Values, and low Crime Rates, low Pupil to Teacher ratio, as well as low % Of Low Status Population.

Question 2 – Data Processing:

a) Summarize the variables in the data set – which variables are continuous, nominal, ordinal?

Continous: Average Daily Balance, Interest Paid, Cash Advances, Balance Transferred, Age Group, Customer Value

Nominal: cust id, Credit Limit, Marital Status, Occupation Group, Customer Type, Gender

Ordinal: LTV Group, Age of Account (Months), Bill Cycle,

b) Identify if there any outliers.

outlier report detailed below.

Quantile Range Outliers

Outliers are values Q times the interquantile range past the lower and upper quantiles.

Tail Quantile 0.1

Q 4

Restrict search to integers

Show only columns with outliers

Rescan

Close

Select columns and choose an action.

Select Rows | Color Cells

Exclude Rows | Color Rows

Add to Missing Value Codes

Change to Missing

Column	10% Quantile	90% Quantile	Low Threshold	High Threshold	Number of Outliers
cust id	30303.9	260212	-889329	1179845	0
Average Daily Balance	63	8622	-34173	42858	21
Interest Paid	11.25	2940	-11704	14655	859
Cash Advances	0	28560.9	-114244	142804	47
Balance Transferred	0	12528.9	-50115	62644.3	692
Age of Account (Months)	18	50	-110	178	0
Age Group	30	46	-34	110	0
Bill Cycle	4	25	-80	109	0
Customer Value	18	7183.98	-28646	35847.9	422
Credit Limit	2000	18000	-62000	82000	0

c) Examine Marital Status, LTV Group, and Gender. Are there any typographical errors? If so, correct them.

For gender, instead of "____", I recoded it to "n/a" for easier reference. In marital status, there were two instances ("N", and "F"). I recoded both to "_____", because it is safer than assuming what they mean, and there is enough data to afford it.

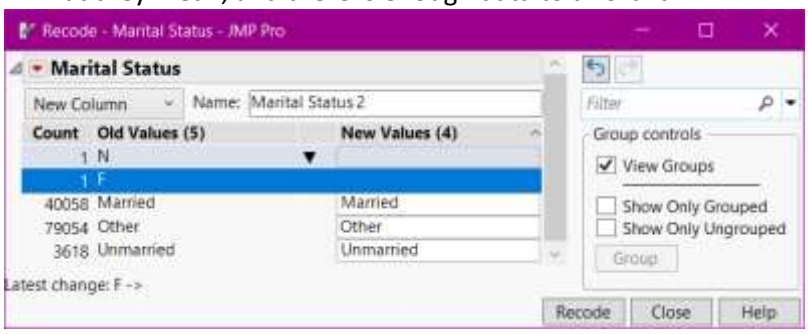

d) Examine Gender. Are there any missing values? If so, how many?

Gender is missing 59,234 entries.

e) Examine Age of Account (Months) and Age Group. Are there any missing values? If so, how many? Impute them

Age of Accounts is missing 68, 798 values, while Age Group is missing 72,912 values. I made up for this with Multivariate Normal Computation.

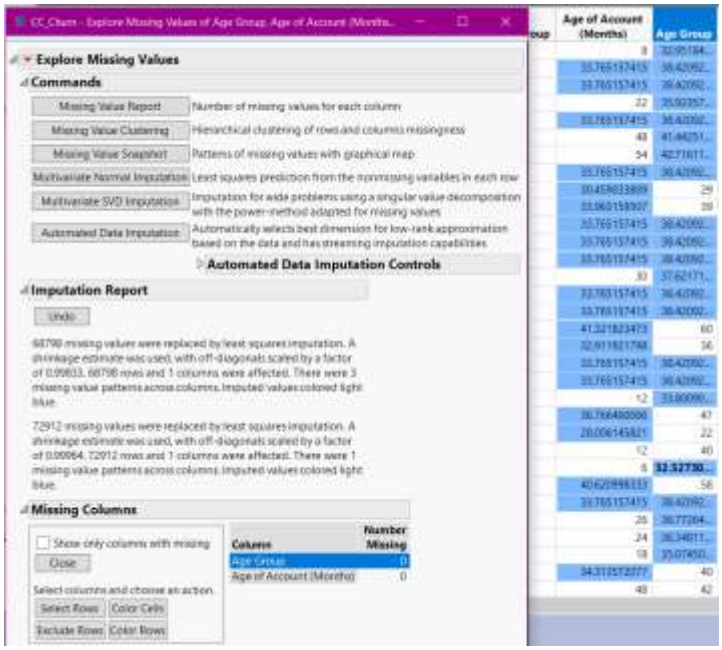

Question 3 – Prediction Using Regression

a. Use the subset that you created in Question 1 as the training data set, and the rest of the data as the validation data set. Compare the variable summaries across the training and validation data sets. Do you have any concerns regarding the partition?

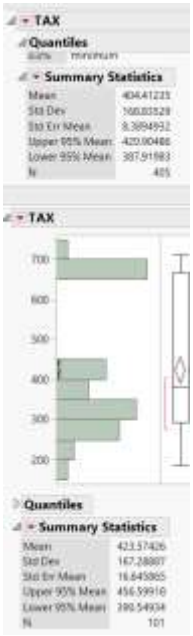

Tax has the most significantly different means, and the most spread (Std. Deviation). The rest of the data does not have this amount of spread.

b. Fit a multiple linear regression model to the median house price (MEDV) as a function of CRIM, CHAS, and RM using the training data Write the equation for predicting the median house price from the predictors in the model.

Parameter Estimates

Term	Estimate	Std Error	t Ratio	Prob>|t|
Intercept	-27.37829	2.918877	-9.38	<.0001*
CRIM	-0.261728	0.036381	-7.19	<.0001*
CHAS[0]	-1.701985	0.604155	-2.82	0.0051*
RM	8.3223961	0.441779	18.84	<.0001*

MEDV = -27.38 -0.26(CRIM) -1.7(CHAS: 0/1) + 8.3(RM)

c. Using the estimated regression model, what median house price is predicted for a tract in the Boston area that does not bound the Charles River, has a crime rate of 0.1, and where the average number of rooms per house is 6? What is the prediction error?
$23,394

d. What is the RMSE on the validation data?
6.3

d. Consider the 12 predictors:

Correlations

	CRIM	ZN	INDUS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	LSTAT	MEDV
CRIM	1.0000	-0.2005	0.4066	0.4210	-0.2192	0.3527	-0.3797	0.6255	0.5828	0.2899	0.4556	-0.3883
ZN	-0.2005	1.0000	-0.5338	-0.5166	0.3120	-0.5695	0.6644	-0.3119	-0.3146	-0.3917	-0.4130	0.3604
INDUS	0.4066	-0.5338	1.0000	0.7637	-0.3917	0.6448	-0.7080	0.5951	0.7208	0.3832	0.6038	-0.4837
NOX	0.4210	-0.5166	0.7637	1.0000	-0.3022	0.7315	-0.7692	0.6114	0.6680	0.1889	0.5909	-0.4273
RM	-0.2192	0.3120	-0.3917	-0.3022	1.0000	-0.2403	0.2052	-0.2098	-0.2920	-0.3555	-0.6138	0.6954
AGE	0.3527	-0.5695	0.6448	0.7315	-0.2403	1.0000	-0.7479	0.4560	0.5065	0.2615	0.6023	-0.3770
DIS	-0.3797	0.6644	-0.7080	-0.7692	0.2052	-0.7479	1.0000	-0.4946	-0.5344	-0.2325	-0.4970	0.2499
RAD	0.6255	-0.3119	0.5951	0.6114	-0.2098	0.4560	-0.4946	1.0000	0.9102	0.4647	0.4887	-0.3816
TAX	0.5828	-0.3146	0.7208	0.6680	-0.2920	0.5065	-0.5344	0.9102	1.0000	0.4609	0.5440	-0.4685
PTRATIO	0.2899	-0.3917	0.3832	0.1889	-0.3555	0.2615	-0.2325	0.4647	0.4609	1.0000	0.3740	-0.5078
LSTAT	0.4556	-0.4130	0.6038	0.5909	-0.6138	0.6023	-0.4970	0.4887	0.5440	0.3740	1.0000	-0.7377
MEDV	-0.3883	0.3604	-0.4837	-0.4273	0.6954	-0.3770	0.2499	-0.3816	-0.4685	-0.5078	-0.7377	1.0000

i. Which predictors are likely to be measuring the same thing among the entire set of predictors? Discuss the relationships among INDUS, NOX, and TAX.

RAD-TAX which have a .91 correlation coefficient. INDUS-NOX have a high coefficient (0.763), as well as INUDS-TAX (0.721), as well as NOX-TAX (0.668).

ii. Compute the correlation table for the numerical predictors and search for highly correlated pairs. These have potential redundancy and can cause multicollinearity. Choose which ones to remove based on this table.

I will remove TAX and NOX, given they are highly correlated with each other and with other predictors.

iii. Use an exhaustive search (All Possible Models) to reduce the remaining predictors as follows: First, choose the top three models. Then run each of these models and compare their predictive accuracy for the validation set. Compare RMSE, Cp, $AICc$, and Validation RSquare. Finally, describe the best model.

All Possible Models

Model	Number	RSquare	RMSE	AICc	BIC	Cp
RM,PTRATIO	2	0.6204	5.5238	2538.77	2554.68	206.6704
RM,PTRATIO,LSTAT	3	0.7265	4.6943	2408.01	2427.88	38.6126
RM,RAD,LSTAT	3	0.6846	5.0409	2465.72	2485.59	105.3268
CRIM,RM,LSTAT	3	0.6828	5.0553	2468.03	2487.90	108.1983
CRIM,RM,PTRATIO,LSTAT	4	0.7327	4.6470	2400.86	2424.68	30.8265
RM,DIS,PTRATIO,LSTAT	4	0.7306	4.6653	2404.04	2427.85	34.1771
RM,RAD,PTRATIO,LSTAT	4	0.7280	4.6872	2407.83	2431.64	38.2143
CRIM,RM,DIS,PTRATIO,LSTAT	5	0.7387	4.6001	2393.70	2421.45	23.2289
INDUS,RM,DIS,PTRATIO,LSTAT	5	0.7369	4.6160	2396.50	2424.25	26.1147
RM,AGE,DIS,PTRATIO,LSTAT	5	0.7354	4.6290	2398.77	2426.51	28.4676
CRIM,INDUS,RM,DIS,PTRATIO,LSTAT	6	0.7442	4.5572	2387.17	2418.84	16.4738
CRIM,RM,AGE,DIS,PTRATIO,LSTAT	6	0.7441	4.5580	2387.31	2418.98	16.6130
CRIM,ZN,RM,DIS,PTRATIO,LSTAT	6	0.7434	4.5647	2388.50	2420.17	17.8111
CRIM,ZN,INDUS,RM,DIS,PTRATIO,LSTAT	7	0.7488	4.5214	2381.85	2417.43	11.0829
CRIM,INDUS,RM,AGE,DIS,PTRATIO,LSTAT	7	0.7485	4.5246	2382.44	2418.01	11.6583
CRIM,ZN,RM,AGE,DIS,PTRATIO,LSTAT	7	0.7475	4.5332	2383.98	2419.56	13.1874
CRIM,ZN,INDUS,RM,AGE,DIS,PTRATIO,LSTAT						
CRIM,ZN,INDUS,RM,DIS,RAD,PTRATIO,LSTAT						
CRIM,INDUS,RM,AGE,DIS,RAD,PTRATIO,LSTAT						
CRIM,ZN,INDUS,RM,AGE,DIS,RAD,PTRATIO,LSTAT						

The best model I found was:

MEDV = CRIM,ZN,INDUS,RM,DIS,PTRATIO,LSTAT, it has a predictive value (Adj. R^2) = 0.744.

Parameter Estimates

Term	Estimate	Std Error	t Ratio	Prob>|t|
Intercept	22.314139	4.240102	5.26	<.0001*
CRIM	-0.119911	0.031723	-3.78	0.0002*
ZN	0.0377B7	0.013953	2.71	0.0071*
INDUS	-0.15319	0.052053	-2.94	0.0034*
RM	4.4368079	0.442731	10.02	<.0001*
DIS	-0.937473	0.187043	-5.01	<.0001*
PTRATIO	-0.877255	0.121574	-7.22	<.0001*
LSTAT	-0.499485	0.048656	-10.27	<.0001*

Effect Tests

Crossvalidation

Source	RSquare	RASE	Freq
Training Set	0.7488	4.4765	405
Validation Set	0.5548	6.7690	101

These are the predictors with their respective estimates, as well as VIF for the Validation Dataset = 6.769.